
AVIS: Autonomous Visual Information Seeking with Large Language Models

Ziniu Hu^{1*} **Ahmet Iscen²** **Chen Sun²** **Kai-Wei Chang¹** **Yizhou Sun¹**

David A Ross²

Cordelia Schmid²

Alireza Fathi²

¹University of California, Los Angeles, ²Google Research

Abstract

In this paper, we propose an autonomous information seeking visual question answering framework, AVIS. Our method leverages a Large Language Model (LLM) to dynamically strategize the utilization of external tools and to investigate their outputs, thereby acquiring the indispensable knowledge needed to provide answers to the posed questions. Responding to visual questions that necessitate external knowledge, such as "What event is commemorated by the building depicted in this image?", is a complex task. This task presents a combinatorial search space that demands a sequence of actions, including invoking APIs, analyzing their responses, and making informed decisions. We conduct a user study to collect a variety of instances of human decision-making when faced with this task. This data is then used to design a system comprised of three components: an LLM-powered planner that dynamically determines which tool to use next, an LLM-powered reasoner that analyzes and extracts key information from the tool outputs, and a working memory component that retains the acquired information throughout the process. The collected user behavior serves as a guide for our system in two key ways. First, we create a transition graph by analyzing the sequence of decisions made by users. This graph delineates distinct states and confines the set of actions available at each state. Second, we use examples of user decision-making to provide our LLM-powered planner and reasoner with relevant contextual instances, enhancing their capacity to make informed decisions. We show that AVIS achieves state-of-the-art results on knowledge-intensive visual question answering benchmarks such as Infoseek [7] and OK-VQA [26].

1 Introduction

Large language models (LLMs), such as GPT3 [5], LaMDA [16], PALM [9], BLOOM [34] and LLaMA [37], have showcased the capacity to memorize and utilize a significant amount of world knowledge. They demonstrate emerging abilities [38] like in-context learning [5], code generation [19], and common sense reasoning [24]. Recently, there is a growing focus towards adapting LLMs to handle multi-modal inputs and outputs involving both vision and language. Noteworthy examples of such visual language models (VLMs) include GPT4 [29], Flamingo [4] and PALI [6]. They set the state of the art for several tasks, including image captioning, visual question answering, and open vocabulary recognition.

While LLMs excel beyond human capabilities in tasks involving textual information retrieval, the current state of the art VLMs perform inadequately on datasets designed for visual information

*This work was done when Ziniu was an intern at Google.

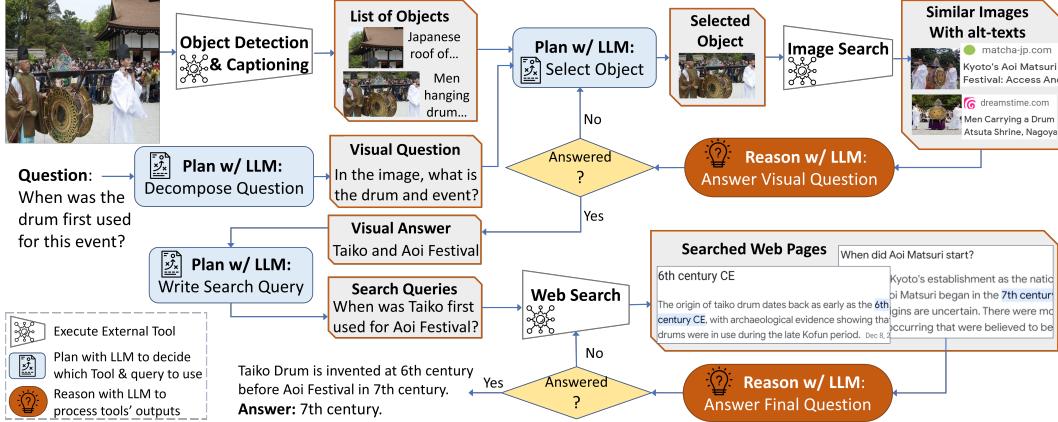


Figure 1: An example of AVIS’s generated workflow for answering a challenging visual question. AVIS consists of an LLM-powered Planner that dynamically selects which tool to use and what query to send, executes the tool, and finally applies an LLM-powered Reasoner processing tool’s outputs.

seeking such as Infoseek [7], Oven [14] and OK-VQA [26]. Many of the visual questions in these datasets are designed in such a way that they pose a challenge even for humans, often requiring the assistance of various APIs and web search to obtain the answer. Examples of such questions include "where is this church located?", "what species of butterfly is this?", or "what is the brand of this dress?".

Current state-of-the-art vision-language models (VLMs) find it challenging to answer such questions for several reasons. Firstly, they are not trained with objectives that encourage them to discern fine-grained categories and details within images. Secondly, they utilize a relatively smaller language model compared to state-of-the-art Large Language Models (LLMs), which constrains their reasoning capabilities. Lastly, they do not compare the query image against a substantial corpus of images associated with varying metadata, unlike systems that employ image search techniques.

To overcome these challenges, we introduce a novel method in this paper that achieves state-of-the-art results on visual information seeking tasks by integrating LLMs with three types of tools: (i) computer vision tools such as object detection, OCR, image captioning models, and VQA models, which aid in extracting visual information from the image, (ii) a web search tool that assists in retrieving open world knowledge and facts, and (iii) an image search tool that enables us to glean relevant information from metadata associated with visually similar images. Our approach utilizes an LLM-powered planner to dynamically determine which tool to use at each step and what query to send to it. Furthermore, we employ an LLM-powered reasoner that scrutinizes the output returned by the tools and extracts the crucial information from them. To retain the information throughout the process, we use a working memory component. Figure 1 shows an example information seeking process performed by our method.

Several recent studies [40, 42, 13, 36, 23] have enhanced LLMs with APIs to handle multi-modal vision-language inputs. These systems generally employ a two-stage strategy, namely *plan* and *execute*. Initially, the LLM breaks down a question into a plan, typically represented as a structured program or a sequence of instructions. Following this, the necessary APIs are activated to collect the required information. While this method has shown potential in elementary visual-language tasks, it frequently fails in more complex real-world situations. In such cases, a comprehensive plan cannot be inferred merely from the initial question. Instead, it necessitates dynamic modifications based on real-time feedback.

The primary innovation in our proposed method lies in its dynamic decision-making capability. Answering visual information seeking questions is a highly complex task, requiring the planner to take multiple steps. At each of these steps, the planner must determine which API to call and what query to send. It is unable to predict the output of complex APIs, such as image search, or to anticipate the usefulness of their responses prior to calling them. Therefore, unlike previous methods that pre-plan the steps and API calls at the beginning of the process, we opt for a dynamic approach.

We make decisions at each step based on the information acquired from previous API calls, enhancing the adaptability and effectiveness of our method.

We conduct a user study to gather a wide range of instances of human decision-making when using APIs to answer questions related to visual information seeking. From this data, we formulate a structured framework that directs the Large Language Model (LLM) to use these examples for making informed decisions regarding API selection and query formulation. The collected user behavior informs our system in two significant ways. First, by analyzing the sequence of user decisions, we construct a transition graph. This graph delineates distinct states and constrains the set of actions available at each state. Second, we use the examples of user decision-making to guide our planner and reasoner with pertinent contextual instances. These contextual examples contribute to improving the performance and effectiveness of our system.

The primary contributions of this paper can be summarized as follows:

- We propose a novel visual question answering framework that leverages a large language model (LLM) to dynamically strategize the utilization of external tools and to investigate their outputs, thereby acquiring the necessary knowledge needed to provide answers to the posed questions.
- We leverage the human decision-making data collected from a user study to develop a structured framework. This framework guides the Large Language Model (LLM) to utilize examples of human decision-making in making informed choices concerning API selection and query construction.
- Our method achieves state-of-the-art results on knowledge-based visual question answering benchmarks such as Infoseek [7] and OK-VQA [26]. Notably, We achieve an accuracy of 50.7% on the Infoseek (unseen entity split) dataset which is significantly higher than the results achieved by PALI [6] with accuracy of 16.0%.

2 Related Work

Augmenting LLMs with Tools. Large Language Models(LLMs) have shown impressive language understanding [33], and even reasoning capabilities [39]. Nevertheless, certain limitations of LLMs are evident, due to their intrinsic characteristics. Such limitations include providing up-to-date answers based on external knowledge or performing mathematical reasoning. Consequently, a recent surge of techniques have integrated LLMs with various external tools [27]. For example, TALM [31] and ToolFormer [35] use in-context learning to teach the language model how to better leverage various tools on benchmarks such as question answering and mathematical reasoning.

In the computer vision domain, LLMs also show significant improvements when combined with external visual tools. For example, Visual ChatGPT [40] and MM-ReAct [42] enable LLMs to call various vision foundation models as tools to understand visual inputs, and even better control the image generation. VisProg [13] and ViperGPT [36] explore the decomposition of visual language tasks into programs, where each line corresponds to general code or a visual API. Chameleon [23] uses an LLM as a natural language planner to infer the appropriate sequence of tools to utilize, and then executes these tools to generate the final response.

Most of these previous works follow a plan-then-execute paradigm, i.e., i) they pre-plan the sequence of actions (API calls) that they will take (either hard coded or using code generation); and ii) they execute the generated plan. One drawback of such an approach is that it cannot update and improve its plan based on the output of the tools it calls. This is not a trivial problem, as it requires to predict the output quality of each tools beforehand. In contrast, our proposed method allows the system to dynamically decide its next steps based on the output it receives from the tools at each step.

Decision Making with LLM as an Agent. There has also been a surge of interest in applying Large Language Models (LLMs) as autonomous agents. These agents are capable of interacting with external environments, making dynamic decisions based on real-time feedback, and consequently achieving specific goals. For example, WebGPT [28] enables an LLM to access real-time information from the web search engines. ReAct [43] further improves external search engine usage via the self-reasoning of LLM in an interleaved manner. Similar ideas have also been adopted for robotic action planning. SayCan [3], for instance, uses LLMs to directly predict robot actions, and PALM-E [10] further fine-tunes LLMs to make better decisions based on instructions and open web media.

When compared to works that follow a plan-then-execute paradigm, these AI agents exhibit increased flexibility, adjusting their actions based on the feedback that they receive. However, many of these methods do not restrict the potential tools that can be invoked at each stage, leading to an immense search space. This becomes particularly critical for web search APIs [1, 2] that return extensive result lists and span a combinatorial search space of multiple tools. Consequently, even the most advanced LLMs today can fall into infinite loops or propagate errors. To alleviate this issue, we propose restricting and guiding LLMs to mimic human behavior when solving complex visual questions with APIs. This idea is similar to the AI alignment research [30, 21] that teaches LLMs to follow human instructions. The difference is that our model only uses the human prior at the decision-making stage via prompt guidance, instead of re-training the model.

3 Method

3.1 General Framework

Our approach employs a dynamic decision-making strategy designed to respond to visual information-seeking queries. Our system is comprised of three primary components. First, we have a planner \mathcal{P} , whose responsibility is to determine the subsequent action, including the appropriate API call and the query it needs to process. Second, we have a working memory \mathcal{M} that retains information about the results obtained from API executions. Lastly, we have a reasoner \mathcal{R} , whose role is to process the outputs from the API calls. It determines whether the obtained information is sufficient to produce the final response, or if additional data retrieval is required.

Considering the potential intricacy of the task, we conduct a user study to gather a broad range of examples of human decision-making process, when using tools to respond to visual information-seeking queries (we introduce the details of data collection in Sec. 3.3). This helps us to establish a structured framework for decision-making. We utilize the data collected from this study to construct a transition graph \mathcal{G} shown in Figure 2, which outlines all the possible actions at each given state. Additionally, we employ real-life decision-making examples \mathcal{E} , i.e., users choose which tool at different states, to guide the planner in choosing the appropriate action at each stage of the process.

The Algorithm 1 presents the operations of the planner \mathcal{P} . The planner undertakes a series of steps each time a decision is required regarding which tool to employ and what query to send to it. Firstly, based on the present *state*, the planner provides a range of potential subsequent actions \mathcal{A}_s . The potential action space \mathcal{A}_s may be large, making the search space intractable. To address this issue, the planner refers to the human decisions from the transition graph \mathcal{G} to eliminate irrelevant actions. The planner also excludes the actions that have already been taken before and are stored in the working memory \mathcal{M} . Formally, this procedure is $\mathcal{A}_s \leftarrow \phi(state, \mathcal{G}, \mathcal{M})$.

Next, it collects a set of relevant in-context examples \mathcal{E}_s that are assembled from the decisions previously made by humans during the user study relevant to actions \mathcal{A}_s , that is $\mathcal{E}_s \leftarrow \theta(\mathcal{E}, \mathcal{A}_s)$. With the gathered in-context examples \mathcal{E}_s and the working memory \mathcal{M} that holds data collected from past tool interactions, the planner formulates a prompt, denoted by $p_s \leftarrow \psi(\mathcal{E}_s, \mathcal{M})$. The prompt p_s is then sent to the LLM which returns a structured answer, determining the next tool t_s to be activated and the query q_s to be dispatched to it. We denote this action by $t_s, q_s \leftarrow LLM(p_s)$. This design allows the planner to be invoked multiple times throughout the process, thereby facilitating dynamic decision-making that gradually leads to answering the input query.

The Algorithm 2 shows the overall decision-making workflow of AVIS. The entire process repeats until a satisfactory answer is produced. Initially, the working memory is populated only with the input visual question I , and the initial *state* is set to START. At each iteration, we first invoke the planner \mathcal{P} to determine the next tool and the query to employ, as outlined in Algorithm 1. Subsequently, the selected external tool executes and delivers its output o_s . The output from the tools can be quite

Algorithm 1 Planner $\mathcal{P}(state, \mathcal{G}, \mathcal{E}, \mathcal{M})$

- 1: $\mathcal{A}_s \leftarrow \phi(state, \mathcal{G}, \mathcal{M})$ \triangleright Get the list of feasible actions \mathcal{A}_s given the current state from transition graph and the information in the working memory
 - 2: $\mathcal{E}_s \leftarrow \theta(\mathcal{E}, \mathcal{A}_s)$ \triangleright Get a list of in-context examples related to actions \mathcal{A}_s
 - 3: $p_s \leftarrow \psi(\mathcal{E}_s, \mathcal{M})$ \triangleright Build a prompt based on the in-context examples \mathcal{E}_s and the current working memory \mathcal{M}
 - 4: $t_s, q_s \leftarrow LLM(p_s)$ \triangleright Decide the next tool t_s to use and the query q_s to pass by feeding the prompt p_s to LLM
-

diverse, ranging from a list of identified objects, to a collection of similar images with their captions, to snippets of search results or knowledge graph entities.

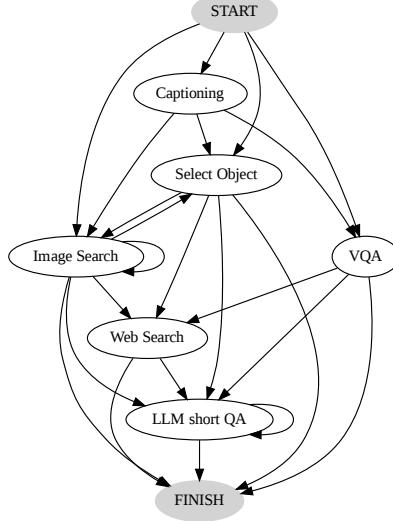


Figure 2: Transition graph \mathcal{G} defines feasible actions the planner can take. This graph is induced by our user study introduced in Sec. 3.3.

Algorithm 2 AVIS Decision Making Workflow

```

1:  $\mathcal{M} \leftarrow \{\text{input}\}$ ,  $\text{state} \leftarrow \text{START}$ 
2:  $t_s, q_s \leftarrow \mathcal{P}(\text{state}, \mathcal{G}, \mathcal{E}, \mathcal{M})$   $\triangleright$  Call the
   planner  $\mathcal{P}$  to decide the next tool to use  $t_s$  and
   the query to pass to it  $q_s$ 
3:  $o_s \leftarrow \text{Exec}(t_s, q_s)$   $\triangleright$  Call tool  $t_s$  with query
    $q_s$  and get output  $o_s$ 
4:  $\hat{o}_s \leftarrow \mathcal{R}(o_s, \mathcal{M})$   $\triangleright$  Process the output and
   extract the key info  $\hat{o}_s$  using the reasoner  $\mathcal{R}$ 
5:  $\mathcal{M}.\text{add}(\hat{o}_s)$   $\triangleright$  Update the working memory
6: switch  $\hat{o}_s$  do
7:   case  $\hat{o}_s$  is not informative
8:     goto(2)  $\triangleright$  Go to line 2 to make
       decision at the same state, excluding  $t_s$ .
9:   case  $\hat{o}_s$  has useful information
10:     $\text{state} \leftarrow t_s$   $\triangleright$  Update state
11:    goto(2)  $\triangleright$  Go to line 2 to make
       decision for the next state.
12:   case  $\hat{o}_s$  is ready as final answer
13:      $ans \leftarrow \hat{o}_s$   $\triangleright$  Output answer
  
```

Therefore, we employ a reasoner \mathcal{R} to analyze the output o_s , extract the useful information and decide into which category the tool output falls: informative, uninformative, or final answer. Our method utilizes the LLM with appropriate prompting and in-context examples to perform the reasoning. If the reasoner concludes that it's ready to provide an answer, it will output the final response, thus concluding the task. If it determines that the tool output is uninformative, it will revert back to the planner to select another action based on the current state. If it finds the tool output to be useful, it will modify the state and transfer control back to the planner to make a new decision at the new state.

To illustrate with a tangible example, we can refer to the output that the model would receive as depicted in Figure 3(c). There are several entities within the answer. The role of the reasoner is twofold: to determine which entity is pertinent for responding to the question and to assess whether the model has obtained the necessary information to transition to the next state.

Our approach, which employs dynamic decision-making coupled with backtracking, differs from previous methods [36, 23] that follow a plan-then-execute paradigm. Our system is structured to make decisions grounded to the results of current executions and to conduct iterative searches for tool combinations. This process eventually yields the most effective strategy to accomplish the task.

3.2 Tools and their APIs

To respond effectively to visual queries that necessitate in-depth information retrieval, it's important to equip AVIS with a comprehensive suite of tools. In this section, we describe these tools.

Image Captioning Model: We employ the PALI 17B [8] captioning model, which obtains state-of-the-art results for image captioning. This tool has the capability to generate captions for either the entire image or for a cropped image corresponding to the bounding box of a detected object.

Visual Question Answering Model: We utilize the PALI 17B [8] VQA model, which has been fine-tuned on the VQA-v2 [11] dataset. This tool takes an image and a question as inputs and provides a text-based answer as the output.

Object Detection: We use an object detector trained on a super-set of Open Images dataset [17] categories that is provided by Google Lens API [1]. We use high confidence threshold to only keep the top-ranked detected boxes for the input image.

Image Search: We utilize Google Image Search to obtain a broad range of information related to the image crop of a detected box as provided in Google Lens API [1]. This information encompasses



Question: In what year was this motorcycle built?



(a) Input visual question and detected objects

PALI Caption
PALI VQA Query: <input type="text"/>
PALI VQA
Search Query: <input type="text"/>
Search
PALM Query: <input type="text"/>
PALM
Show entity of box 1
Show caption of images similar to box 1
Show entity of box 2 (highlighted)
Show related products to the object in box 2
Show caption of images similar to box 2
Show caption of identical images to box 2
Show entity of box 3
Show caption of images similar to box 3
Show related products to the object in box 4
Show caption of images similar to box 4
Show caption of images similar to box 5

(b) Tools shown to user

Outputs of “show entity of box2”

- Harley-Davidson
- Harley-Davidson, Inc. is an American motorcycle manufacturers to survive the Great Depression & arrangements, periods of poor economic health brand widely known for its loyal following. There
- Harley-Davidson is noted for a style of customiz motorcycles with engine displacements greater
- Harley-Davidson manufactures its motorcycles :
- Despatch rider
- A despatch rider is a military messenger, mount
- In the UK ‘despatch rider’ is also a term used for
- Despatch riders were used by armed forces to d telecommunications were limited and insecure.

Useless API call

Success! Found the Answer!

Couldn't find the Answer!

(c) Tool Output

Figure 3: We conduct a user study to gather examples of user decision-making when responding to visual information-seeking questions. Given a visual question as depicted in (a), the user makes a series of tool calls using the available APIs shown in (b). Each tool call yields an output which the user reviews whether it is useful and determines the subsequent action, illustrated in (c).

various details, such as knowledge graph entities, titles of associated products, and captions of analogous or identical images. The availability of these details can vary based on the image crop input provided to Google Image Search. When it comes to decision-making, our planner considers the utilization of each piece of information as a separate action. This is due to the fact that each information could contain hundreds of tokens that necessitate complex processing and reasoning.

OCR: In some cases, images may include textual content such as street names or logos. To detect and utilize this text, we take advantage of the Optical Character Recognition (OCR) feature available in the Google Lens API [1].

Web Search: Web search enables our approach to acquire up-to-date world knowledge and retrieve relevant documents on any topic of interest. For this objective, we employ the Google Web Search API [2]. It accepts a text-based query as input and produces the following outputs: (i) related document links and snippets, (ii) in certain instances, a knowledge panel providing a direct answer to the query, and (iii) up to five questions that are related to the input query. If a knowledge panel is available, we parse it into a sentence or a few sentences that summarize its information.

LLM short QA: We incorporate a Language Model (LLM) powered question-answering component as another tool. This tool accepts a query in text form and produces an answer also in text form. It is important to note that the use of the LLM here as a question-answering tool is distinct from its role in the planner or reasoner as outlined in Alg. 1 and Alg. 2.

3.3 Gathering User Behavior to Inform LLM Decision Making

Many of the visual questions in datasets such as Infoseek [7], Oven [14] and OK-VQA [26] ask for fine-grained answers, which poses a challenge even for humans, often requiring the assistance of various APIs and web searches for answers. Figure 3(a) illustrates an example visual question taken from the OK-VQA [26] dataset. In order to gather insights into human decision-making process, we carried out a user study. More specifically, our goal is to understand how humans utilize external tools to answer visual queries that involve seeking information.

The user is equipped with an identical set of tools as our method. They are presented with the input image and question, along with image crops for each detected object. Additionally, tools like PALI Caption, PALI VQA, PALM, and Web Search are made available to the user. Furthermore, based on the information obtained through image search for each cropped image, the user is offered one or multiple buttons associated with each box. These buttons provide the user with the ability to access diverse information pertaining to the image crop of the box. This includes details such as corresponding knowledge graph entities, captions of similar images, titles of associated related products, and captions of identical images. An example set of tools and APIs are shown in Figure 3(b).

When the user initiates an action, such as clicking on a button or submitting a query to web search, PALM, or PALI VQA, the corresponding tool is invoked, and the resulting output is displayed to the user. We record the sequence of actions taken by the user and the outputs that they receive at each step. For instance, in Figure 3, we show an example of how a user needs to perform four actions to answer the question: *i*) display entities in box 2, *ii*) show the caption of similar images to box 2, *iii*) conduct a search for "*In what year was Harley-Davidson XA built?*", and *iv*) utilize PALM using the combination of the search output and the question "*In what year was Harley-Davidson XA built?*". When the user is prepared to proceed to the next question, they click on either of the two buttons: "Success! Found the Answer!" or "Couldn't Find the Answer." Subsequently, a new visual question is presented to them.

The collected user behavior serves as a guide for our system in two key ways. Firstly, we construct a transition graph by analyzing the sequence of decisions made by users. This graph defines distinct states and restricts the available set of actions at each state. For example, at the START state, the system can take only one of these three actions: PALI caption, PALI VQA, or object detection. Figure 2 illustrates the transition graph that has been constructed based on the decision-making process of the users. Secondly, we utilize the examples of user decision-making to guide our planner and reasoner with relevant contextual instances. These in-context examples aid in enhancing the performance and effectiveness of our system.

We conducted a user study involving 10 participants who collectively answered a total of 644 visual questions. During the study, we presented users with visual questions that were randomly selected from both the Infoseek [7] and OK-VQA [26] datasets. This approach allowed us to provide the participants with a varied and diverse set of visual questions to assess and respond to. We show the details for this study as well as example prompts in the Appendix.

4 Experiments

We evaluate AVIS on two visual question answering datasets: *i*) OK-VQA [26], which requires common-sense knowledge not observed in given image; and *ii*) Infoseek_{wikidata} [7], which further necessitates more fine-grained information that cannot be covered by common sense knowledge.

Experimental Setup. We follow the decision-making workflow in Alg. 2 to implement AVIS to solve visual questions. For the Planner, we write the basic instructions for describing each tool, and keep a pool of real user behavior when they select each tool, which we collected in the user study. At each step s , we prepare the prompt based on the feasible action lists \mathcal{A}_s . For the Reasoner, we write the prompt for all APIs that return a long list of results, including *Object Detection*, *Product Detection*, *Web Image Search* and *Web Text Search*, that guides reasoner to extract the relevant information. Note that we design the reasoner in a way such that the “uninformative” answers can be detected. In order to support this, we manually prepare several bad examples that do not provide any useful information, pass it to the reasoner as a part of the prompt. We show the detailed prompts for these two modules in the Appendix.

We use the frozen PALM 540B language model [9] for both the planner and the reasoner, with deterministic generation ensured by setting the temperature parameter to zero. We use 10 examples as in-context prompts for each dataset, and report the VQA accuracy [11] as the evaluation metric.

Baselines. A significant novelty of AVIS is the ability to dynamically determine the relevant tools according to different states. To show that this design choice is useful, we add a number of baselines that do not contain a LLM-planner for dynamic decision making. Instead, they follow a pre-determined sequence to call a list of tools. We propose the following baselines:

- **baseline-PALM w/ PALI***, which integrates the captions generated by PALI and the visual answers from PALI VQA. PALI* denotes the combination of both VQA and captioning tool.
- **baseline-PALM w/ (PALI* + Object)**, which in addition calls the object detection tool, and then integrates all object data, including products and text detected by OCR.
- **baseline-PALM w/ (PALI* + Object + Search)**, a model which first selects a relevant object with the help of PALM, then sequentially executes the image search and Google search with the object name. It then calls PALM again to answer the question.

For each of the three baselines, we prepare a few-shot Chain-Of-Thought (COT) prompting [39], in which the COT prompt guides the model to explain why predictions are made based on the provided

Model	Unseen Entity	Unseen Question
PALM [9] (Q-only, few-shot)	3.7	5.1
OFA [22] (fine-tune)	9.7	14.8
PALI [6] (VQA, zero-shot)	1.8	2.2
PALI [6] (fine-tune)	16.0	<u>20.7</u>
PALM [9] w/ CLIP [32] (few-shot + external knowledge)	<u>21.9</u>	18.6
FiD [44] w/ CLIP [32] (fine-tune + external knowledge)	20.7	18.1
(—baselines without dynamic decision making, sequentially execute the tools—)		
baseline-PALM w/ (PALI*, few-shot)	12.8	14.9
baseline-PALM w/ (PALI* + Object, few-shot)	31.3	36.1
baseline-PALM w/ (PALI* + Object + Search, few-shot)	36.1	38.2
AVIS (ours, few-shot)	50.7	56.4
w/o PALI*	47.9	54.2
w/o Object	41.2	48.4
w/o Search	42.5	49.6

Table 1: **Visual Question Answering** results (accuracy) on Infoseek_{Wikidata}. The first four rows are results from their paper that do not use external knowledge, and the next two are from their paper that use CLIP as knowledge source. The tool PALI* denotes the frozen multi-task PALI-17B model for both visual question answering and image captioning. Object means object detection, and search means image and text search.

Model	Accuracy (%)
Supervised	KRISP [25]
	KAT [12]
	ReViVE [20]
	REVEAL [15]
	PALI [6] (OK-VQA, finetune)
Zero-shot	PALI [6] (VQA, zero-shot)
	PICa-Full [41]
	Flamingo (zero-shot) [4]
	BLIP-2 [18]
	ViperGPT [36]
Few-shot	Flamingo (few-shot) [4]
	(baselines without dynamic decision making, sequentially executing the tools)
	baseline-PALM w/ (PALI*)
	baseline-PALM w/ (PALI*+Object)
	baseline-PALM w/ (PALI*+Object + Search)
AVIS (ours)	60.2
w/o PALI*	47.1
w/o Object	58.3
w/o Search	55.0

Table 2: **Visual Question Answering** results (accuracy) on OK-VQA. The tool PALI* denotes the frozen multi-task PALI-17B model for both visual question answering and image captioning. Object means object detection, and search means image and text search.

information. Note that these baselines utilize a set of tools in a fixed order, without the capacity for dynamic decision making.

We also evaluate the usefulness of each tool group (i.e., PALI*, Object, and Search) through an ablation study. This involves removing each tool group from our framework individually, and assessing the impact on performance.

Experimental Results. Table 5 presents the results of AVIS and other baselines on the Infoseek_{Wikidata} dataset. Infoseek_{Wikidata} is a challenging dataset that requires identifying highly specific entities. Even robust visual-language models, such as OFA [22] and PALI [6], fail to yield

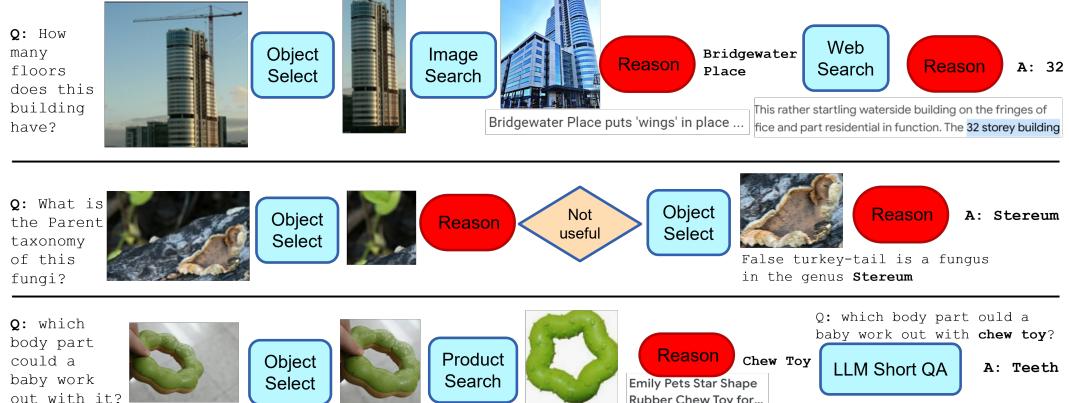


Figure 4: Examples of AVIS’s dynamic planning and reasoning procedure for solving visual questions.

high accuracy when fine-tuned on this dataset. However, our AVIS, without fine-tuning and by leveraging a complete set of tools guided by 10 in-context examples, achieves the accuracy of 50.7 and 56.4 on the unseen entity and question splits, respectively. This significantly outperforms the fine-tuned results of PALI-17B, which are 16.0 and 20.7, as well as the PALM model augmented with CLIP knowledge, which are 21.9 and 18.6, respectively.

Table 5 also illustrates that our improvements are not solely due to the additional information provided by the external tools, but due to our dynamic decision-making pipeline. We compare the results of AVIS with the three baselines that conduct sequential execution. While these baselines do improve the performance, our AVIS framework outperforms the best baseline model by up to 17.3 accuracy. Note that AVIS and the baselines use exactly the same set of tools. This considerable performance gap clearly shows the clear advantage of our dynamic decision-making design. Furthermore, we show the importance of each tool in the last block of Table 5. Removal of any of the tools degrades the overall accuracy. Among the three tool groups, Object and Search are more important than PALI, as they provide more fine-grained information crucial for the Infoseek dataset.

We report the OK-VQA experiments in Table 2. AVIS with few-shot in-context examples achieves an accuracy of 60.2, higher than most of the existing methods tailored for this dataset, including KAT [12], ReVIVE [20] and REVEAL [15]. AVIS achieves lower but comparable performance compared to PALI model fine-tuned on OK-VQA. This difference, compared to Infoseek, may be attributed to the fact that most QA examples in OK-VQA rely more on commonsense knowledge than on fine-grained knowledge. Therefore, it is feasible to encode such generic knowledge in the model parameters and requires less external knowledge. Note that PALI zero-shot VQA model itself achieves 41.6 accuracy, which is significantly higher than in Infoseek, which supports this hypothesis. Table 2 also shows that the object detection is less crucial as a tool on this data set, compared to PALI captioning and VQA. Nonetheless, AVIS equipped with all tools achieves the best performance.

Case studies for dynamic decision making. One of the key features of AVIS is its ability to dynamically make decisions instead of executing a fixed sequence. Figure 4 presents three examples of AVIS’s dynamic planning and reasoning process. They demonstrate the flexibility of AVIS to use different tools at various stages. It is also worth noting that our reasoner design enables AVIS to identify irrelevant information, backtrack to a previous state, and repeat the search. For instance, in the second example concerning the taxonomy of fungi, AVIS initially makes an incorrect decision by selecting a leaf object. However, the reasoner identifies that this is not relevant to the question, prompting AVIS to plan again. This time, it successfully selects the object related to false turkey-tail fungi, leading to the correct answer, Stereum.

5 Conclusion

In this paper, we propose a novel approach that equips the Large Language Models (LLM) with the ability to use a variety of tools for answering knowledge-intensive visual questions. Our methodology, anchored in human decision-making data collected from a user study, employs a structured framework that uses an LLM-powered planner to dynamically decide on tool selection and query formation. An

LLM-powered reasoner is tasked with processing and extracting key information from the output of the selected tool. Our method iteratively employs the planner and reasoner to leverage different tools until all necessary information required to answer the visual question is amassed.

Limitation Statement: Currently AVIS is specifically designed for visual question answering. We aim to extend our LLM-powered dynamic decision-making framework to address other reasoning tasks. Additionally, our current framework depends on a computationally intensive LLM, namely, the PALM model. We are interested in investigating whether this decision-making framework can also be performed by lighter weight language models.

References

- [1] Google lens, 2023. Web interface available at <https://images.google.com>.
- [2] Google search, 2023. Web interface available at <https://www.google.com>.
- [3] Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Keerthana Gopalakrishnan, Karol Hausman, Alex Herzog, et al. Do as i can, not as i say: Grounding language in robotic affordances. *arXiv preprint arXiv:2204.01691*, 2022.
- [4] Jean-Baptiste Alayrac, Jeff Donahue, Pauline Luc, Antoine Miech, Iain Barr, Yana Hasson, Karel Lenc, Arthur Mensch, Katherine Millican, Malcolm Reynolds, et al. Flamingo: a visual language model for few-shot learning. *Advances in Neural Information Processing Systems*, 35:23716–23736, 2022.
- [5] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- [6] Xi Chen, Xiao Wang, Soravit Changpinyo, AJ Piergiovanni, Piotr Padlewski, Daniel Salz, Sebastian Goodman, Adam Grycner, Basil Mustafa, Lucas Beyer, et al. Pali: A jointly-scaled multilingual language-image model. *arXiv preprint arXiv:2209.06794*, 2022.
- [7] Yang Chen, Hexiang Hu, Yi Luan, Haitian Sun, Soravit Changpinyo, Alan Ritter, and Ming-Wei Chang. Can pre-trained vision and language models answer visual information-seeking questions? In *arXiv preprint arXiv:2302.11713*, 2023.
- [8] Felix Chern, Blake Hechtman, Andy Davis, Ruiqi Guo, David Majnemer, and Sanjiv Kumar. TPU-KNN: K nearest neighbor search at peak flop/s. *CoRR*, abs/2206.14286, 2022.
- [9] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. Palm: Scaling language modeling with pathways. *arXiv preprint arXiv:2204.02311*, 2022.
- [10] Danny Driess, Fei Xia, Mehdi SM Sajjadi, Corey Lynch, Aakanksha Chowdhery, Brian Ichter, Ayzaan Wahid, Jonathan Tompson, Quan Vuong, Tianhe Yu, et al. Palm-e: An embodied multimodal language model. *arXiv preprint arXiv:2303.03378*, 2023.
- [11] Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. Making the V in VQA matter: Elevating the role of image understanding in visual question answering. In *2017 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2017, Honolulu, HI, USA, July 21-26, 2017*, pages 6325–6334. IEEE Computer Society, 2017.
- [12] Liangke Gui, Borui Wang, Qiuyuan Huang, Alex Hauptmann, Yonatan Bisk, and Jianfeng Gao. Kat: A knowledge augmented transformer for vision-and-language. *arXiv preprint arXiv:2112.08614*, 2021.
- [13] Tanmay Gupta and Aniruddha Kembhavi. Visual programming: Compositional visual reasoning without training. In *arXiv preprint arXiv:2211.11559*, 2022.
- [14] Hexiang Hu, Yi Luan, Yang Chen, Urvashi Khandelwal, Mandar Joshi, Kenton Lee, Kristina Toutanova, and Ming-Wei Chang. Open-domain visual entity recognition: Towards recognizing millions of wikipedia entities. In *arXiv preprint arXiv:2302.11154*, 2023.
- [15] Ziniu Hu, Ahmet Iscen, Chen Sun, Zirui Wang, Kai-Wei Chang, Yizhou Sun, Cordelia Schmid, David A. Ross, and Alireza Fathi. Reveal: Retrieval-augmented visual-language pre-training with multi-source multimodal knowledge memory. In *CVPR*, 2023.
- [16] Apoorv Kulshreshtha, Daniel De Freitas Adiwardana, David Richard So, Gaurav Nemade, Jamie Hall, Noah Fiedel, Quoc V. Le, Romal Thoppilan, Thang Luong, Yifeng Lu, and Zi Yang. Towards a human-like open-domain chatbot. In *arXiv preprint arXiv:2001.09977*, 2020.
- [17] Alina Kuznetsova, Hassan Rom, Neil Alldrin, Jasper Uijlings, Ivan Krasin, Jordi Pont-Tuset, Shahab Kamali, Stefan Popov, Matteo Malloci, Alexander Kolesnikov, Tom Duerig, and Vittorio Ferrari. The open images dataset v4: Unified image classification, object detection, and visual relationship detection at scale. *IJCV*, 2020.

- [18] Junnan Li, Dongxu Li, Silvio Savarese, and Steven C. H. Hoi. BLIP-2: bootstrapping language-image pre-training with frozen image encoders and large language models. *CoRR*, abs/2301.12597, 2023.
- [19] Yujia Li, David H. Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, Thomas Hubert, Peter Choy, Cyprien de Masson d’Autume, Igor Babuschkin, Xinyun Chen, Po-Sen Huang, Johannes Welbl, Sven Gowal, Alexey Cherepanov, James Molloy, Daniel J. Mankowitz, Esme Sutherland Robson, Pushmeet Kohli, Nando de Freitas, Koray Kavukcuoglu, and Oriol Vinyals. Competition-level code generation with alphacode. *CoRR*, abs/2203.07814, 2022.
- [20] Yuanze Lin, Yujia Xie, Dongdong Chen, Yichong Xu, Chenguang Zhu, and Lu Yuan. Revive: Regional visual representation matters in knowledge-based visual question answering. *arXiv preprint arXiv:2206.01201*, 2022.
- [21] Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *arXiv preprint arXiv:2304.08485*, 2023.
- [22] Jiasen Lu, Christopher Clark, Rowan Zellers, Roozbeh Mottaghi, and Aniruddha Kembhavi. Unified-io: A unified model for vision, language, and multi-modal tasks. *CoRR*, abs/2206.08916, 2022.
- [23] Pan Lu, Baolin Peng, Hao Cheng, Michel Galley, Kai-Wei Chang, Ying Nian Wu, Song-Chun Zhu, and Jianfeng Gao. Chameleon: Plug-and-play compositional reasoning with large language models. In *arXiv preprint arXiv:2304.09842*, 2023.
- [24] Aman Madaan, Shuyan Zhou, Uri Alon, Yiming Yang, and Graham Neubig. Language models of code are few-shot commonsense learners. In Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang, editors, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 1384–1403. Association for Computational Linguistics, 2022.
- [25] Kenneth Marino, Xinlei Chen, Devi Parikh, Abhinav Gupta, and Marcus Rohrbach. Krisp: Integrating implicit and symbolic knowledge for open-domain knowledge-based vqa. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14111–14121, 2021.
- [26] Kenneth Marino, Mohammad Rastegari, Ali Farhadi, and Roozbeh Mottaghi. OK-VQA: A visual question answering benchmark requiring external knowledge. In *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2019, Long Beach, CA, USA, June 16-20, 2019*, pages 3195–3204. Computer Vision Foundation / IEEE, 2019.
- [27] Grégoire Mialon, Roberto Dessì, Maria Lomeli, Christoforos Nalmpantis, Ram Pasunuru, Roberta Raileanu, Baptiste Rozière, Timo Schick, Jane Dwivedi-Yu, Asli Celikyilmaz, Edouard Grave, Yann LeCun, and Thomas Scialom. Augmented language models: a survey. In *arXiv preprint arXiv:2302.07842*, 2023.
- [28] Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, et al. Webgpt: Browser-assisted question-answering with human feedback. *arXiv preprint arXiv:2112.09332*, 2021.
- [29] OpenAI. Gpt-4 technical report. In *arXiv preprint arXiv:2303.08774*, 2023.
- [30] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.
- [31] Aaron Parisi, Yao Zhao, and Noah Fiedel. Talm: Tool augmented language models. In *arXiv preprint arXiv:2205.12255*, 2022.
- [32] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event*, volume 139 of *Proceedings of Machine Learning Research*, pages 8748–8763. PMLR, 2021.
- [33] Alec Radford, Karthik Narasimhan, Tim Salimans, Ilya Sutskever, et al. Improving language understanding by generative pre-training. 2018.
- [34] Teven Le Scao, Angela Fan, Christopher Akiki, Ellie Pavlick, Suzana Ilic, Daniel Hesslow, Roman Castagné, Alexandra Sasha Luccioni, François Yvon, Matthias Gallé, Jonathan Tow, Alexander M. Rush, Stella Biderman, Albert Webson, Pawan Sasanka Ammanamanchi, Thomas Wang, Benoit Sagot, Niklas Muennighoff, Albert Villanova del Moral, Olatunji Ruwase, Rachel Bawden, Stas Bekman, Angelina McMillan-Major, Iz Beltagy, Huu Nguyen, Lucile Saulnier, Samson Tan, Pedro Ortiz Suarez, Victor Sanh, Hugo Laurençon, Yacine Jernite, Julien Launay, Margaret Mitchell, Colin Raffel, Aaron Gokaslan, Adi Simhi, Aitor Soroa, Alham Fikri Aji, Amit Alfassy, Anna Rogers, Ariel Kreisberg Nitzav, Canwen Xu, Chenghao Mou, Chris Emezue, Christopher Klamm, Colin Leong, Daniel van Strien, David Ifeoluwa Adelani, and et al. BLOOM: A 176b-parameter open-access multilingual language model. *CoRR*, abs/2211.05100, 2022.

- [35] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *arXiv preprint arXiv:2302.04761*, 2023.
- [36] Dídac Surís, Sachit Menon, and Carl Vondrick. Vipergpt: Visual inference via python execution for reasoning. In *arXiv preprint arXiv:2303.08128*, 2023.
- [37] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- [38] Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, Ed H. Chi, Tatsunori Hashimoto, Oriol Vinyals, Percy Liang, Jeff Dean, and William Fedus. Emergent abilities of large language models. In *arXiv preprint arXiv:2206.07682*, 2022.
- [39] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *NeurIPS*, 2022.
- [40] Chenfei Wu, Shengming Yin, Weizhen Qi, Xiaodong Wang, Zecheng Tang, and Nan Duan. Visual chatgpt: Talking, drawing and editing with visual foundation models. In *arXiv preprint arXiv:2303.04671*, 2023.
- [41] Zhengyuan Yang, Zhe Gan, Jianfeng Wang, Xiaowei Hu, Yumao Lu, Zicheng Liu, and Lijuan Wang. An empirical study of GPT-3 for few-shot knowledge-based VQA. *ArXiv preprint, abs/2109.05014*, 2021.
- [42] Zhengyuan Yang, Linjie Li, Jianfeng Wang, Kevin Lin, Ehsan Azarnasab, Faisal Ahmed, Zicheng Liu, Ce Liu, Michael Zeng, and Lijuan Wang. Mm-react: Prompting chatgpt for multimodal reasoning and action. In *arXiv preprint arXiv:2303.11381*, 2023.
- [43] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*, 2022.
- [44] Donghan Yu, Chenguang Zhu, Yuwei Fang, Wenhao Yu, Shuohang Wang, Yichong Xu, Xiang Ren, Yiming Yang, and Michael Zeng. KG-FiD: Infusing knowledge graph in fusion-in-decoder for open-domain question answering. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4961–4974, Dublin, Ireland, 2022. Association for Computational Linguistics.

A Implementation of AVIS workflow

We implemented AVIS using the code snippet referenced in Code 7. Throughout our experiments, we employed the APIs of Google Search, LENS, PALI, and PALM directly, without the need for additional GPU/TPU computational resources. Tools that didn't require input queries, such as object detection, captioning, and image search, had their results pre-calculated over the two datasets to reduce the time cost. Other services like VQA, text search, and LLM QA were called during runtime.

B Comparison to pure Autonomous baseline without Transition Graph

One of the significant contributions of this paper lies in the use of a transition graph, synthesized from an authentic user study. To underscore the importance of this graph, along with user prompts in facilitating the efficacy of AVIS, we devised a baseline that operates independently of the transition graph. In this scenario, the model, at each timestep, is presented with a comprehensive list of all tools, each paired with a task description. This baseline shares similarities with the recently launched AutoGPT², BabyAGI³ projects, which attempted to utilize LLMs as autonomous agents to select all possible actions available in the web.

The results are show in Table 3 on Infoseek WIkiData unseen entity set and OKVQA. Note that this baseline doesn't achieve the number as high as AVIS with the transition graph and user prompts. The key reason for this discrepancy is the global characteristics inherent in the tool list we have. For instance, we typically first address the visual sub-question through object detection and image search, followed by resolving the knowledge component via Google Search and LLM. However, solely relying on the task description, devoid of human behavior as guidance, can result in the model generating unrealistic tools. We will discuss this intuition more in the following sections.

Model	Infoseek	OKVQA
AVIS w/o/ Transition Graph	38.2	47.3
AVIS w/ Transition Graph	50.7	60.2

Table 3: Ablation of AVIS with or without the guidance of Transition Graph

C Analysis of AVIS's generated tool execution sequence

We have also conducted an analysis to determine whether common patterns exist within the generated programs of AVIS's predictions.

We gathered the tool execution traces for all samples within the Infoseek unseen entity dataset. Initially, we display the frequency of each tool being invoked in Figure 5, followed by a more detailed analysis of the first to fourth most commonly called tools in Figures 6-9. As illustrated, the AVIS model, guided by the transition graph and prompts, does not utilize all possible combination of tools, but favors some certain combinations. For instance, as depicted in Fig 6, "object select" is utilized more frequently than other tools at the outset. Similarly, as demonstrated in Fig 8, during the third step, when the model accumulates the visual answer, it is likely to invoke "web search" to gather additional information.

We have also calculated the transition probability of the induced graph in Fig 10. The structure of this graph differs slightly from the guided transition graph because during actual runtime, the model will not predict some of the edges. Overall, it reveals a clear two-step question-solving pattern. Initially, AVIS gathers sufficient visual information through the use of visual tools such as "object detection," "VQA," or "identical image search," and then employs "LLM QA" to obtain the visual answer. Subsequently, it iteratively calls "web search" and "LLM QA" post-search with a prompt, eventually deriving the final answer. We also present the distribution of the lengths of generated sequences in Figure 12. As illustrated, the lengths vary considerably, rather than maintaining a fixed value, with a length of 5 being most common for the generated sequences.

²<https://github.com/Significant-Gravitas/Auto-GPT>

³<https://github.com/yoheinakajima/babyagi>

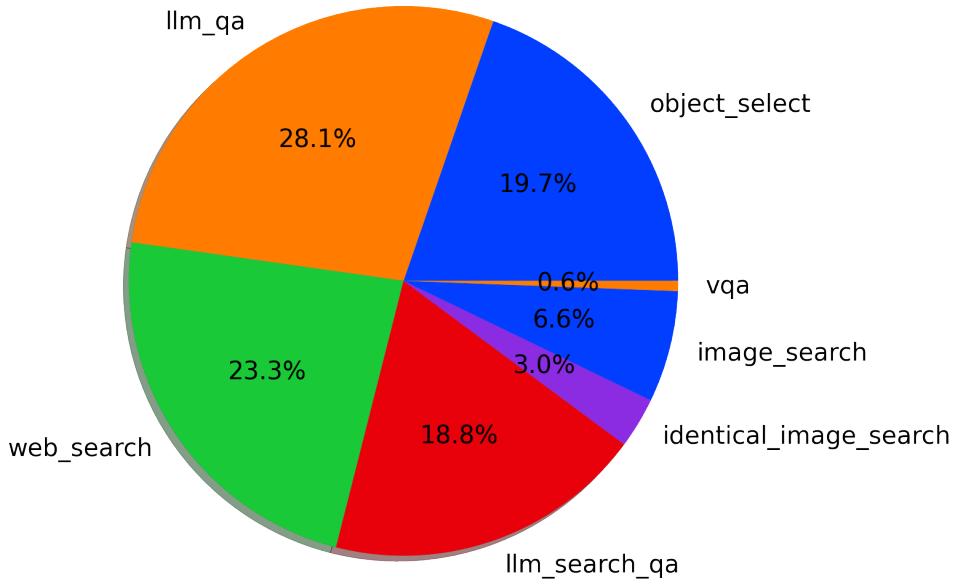


Figure 5: Overall frequency of tool usage on Infoseek dataset.

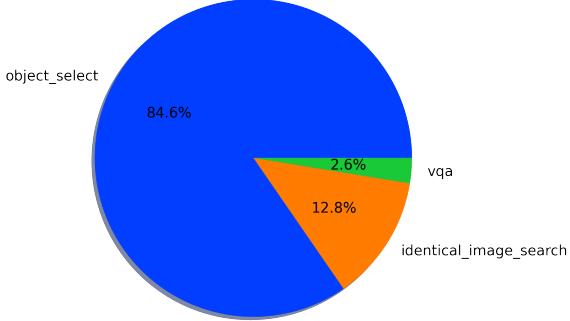


Figure 6: Frequency of the first used tool.

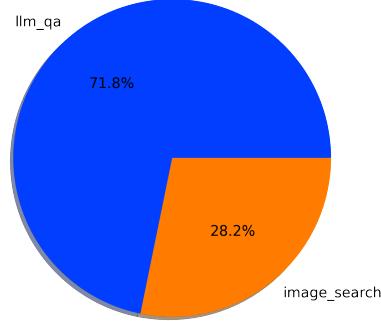


Figure 7: Frequency of the second used tool.

Another intriguing aspect worth exploring is our reasoner component. As explained in the paper, the reasoner evaluates whether the output of each tool is "informative," "not informative," or "answerable". We exhibit the overall frequency of these predictions in Figure 11. As shown, the model tends to classify most of the outputs as either informative or answerable. However, approximately 8.1% of returned entries are deemed "not informative," in which case AVIS would backtrack to select alternative actions. We further demonstrate a few examples of different choices in Table 4.

D Dataset Details

Infoseek⁴ is a Visual Question Answering (VQA) dataset, specifically geared towards information-seeking questions that cannot be answered merely through common sense knowledge. This dataset was curated by initially gathering human-annotated questions, which were then automatically integrated with existing visual entity recognition datasets and Wikidata to generate complex question-

⁴<https://open-vision-language.github.io/infoseek/>

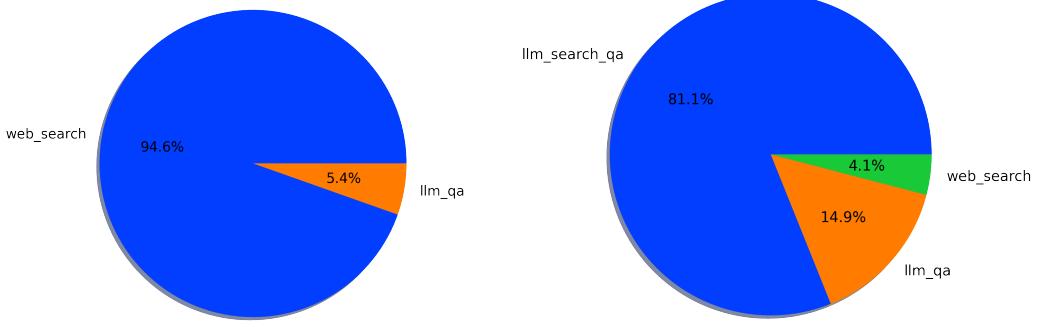


Figure 8: Frequency of the third used tool.

Figure 9: Frequency of the forth used tool.

Question	Action	Returned Content	Output of Reasoner
where is the island located in?	image search	Retrievals: [willow - Students Britannica Kids Homework Help (score=57.2), Silky Willow Live Stakes For Sale Buy Live Stakes Wholesale (score=55.3), ...]	not informative
Who or what maintains this place?	LLM QA	the history of trafalgar square. the most relevant knowledge entry about trafalgar square is "Trafalgar Square is the focal point for protests, revelry and victory celebrations, however it was only in 1835 that the area became known by this name". From the result, we know the history of trafalgar square is trafalgar square is the focal point for protests, revelry and victory celebrations, however it was only in 1835 that the area became known by this name. Terefore, the predicted answer is trafalgar square is the focal point for protests , revelry and victory celebrations , however it was only in 1835 that the area became known by this name.	not informative
What is the total quantity of produced items for this type of aircraft?	image search	Retrievals: [high confidence: Condor bietet neue Langstrecke nach Brasilien an - AERO International (score=97.3), high confidence: Air safety incidents for aircraft registration D-ABOE - AeroInside (score=95.0), ...]	yes, answerable
what is the name of this mountain?	image search	Object: [a view of a castle in the distance . There are trees and buildings present at the bottom of this image. We can see a hill and a tower in the middle of this image. We can see the sky in the background. (Caption, whole image). Mount of Olives (ridge): The Mount of Olives or Mount Olivet is a mountain ridge east of and adjacent to Jerusalem's Old City (score=88.6), Mount Zion (peak): Mount Zion is a hill in Jerusalem, located just outside the walls of the Old City (score=79.0)]	yes, informative

Table 4: Several examples of API execution results and the reasoner’s justification.

answer pairs. At the time of submission, we only have access to its wikidata split. Here we also report the results on human split in Table 5.

OK-VQA⁵ is another VQA dataset, unique in its requirement for the application of external knowledge that transcends the information directly visible in the input images. The creation of this dataset involved crowdsourced workers who were tasked with annotating complex questions, drawing upon the extensive knowledge resources available on Wikipedia.

E Prompt Examples

Below we show different prompt examples to support our AVIS workflow. First is the prompts for planning, which selects which tool to use and what query to send. It is consists of a overall task descriptions and many real examples showing at which circumstances real users select this tool.

```

1planner_prompt =
2"""You goal is to answer the following query: %s.
3
4To answer it, you will be provided with the following tools:
5%s
6
7Please make the decision based on the current context.
8
9%s

```

⁵<https://okvqa.allenai.org/>

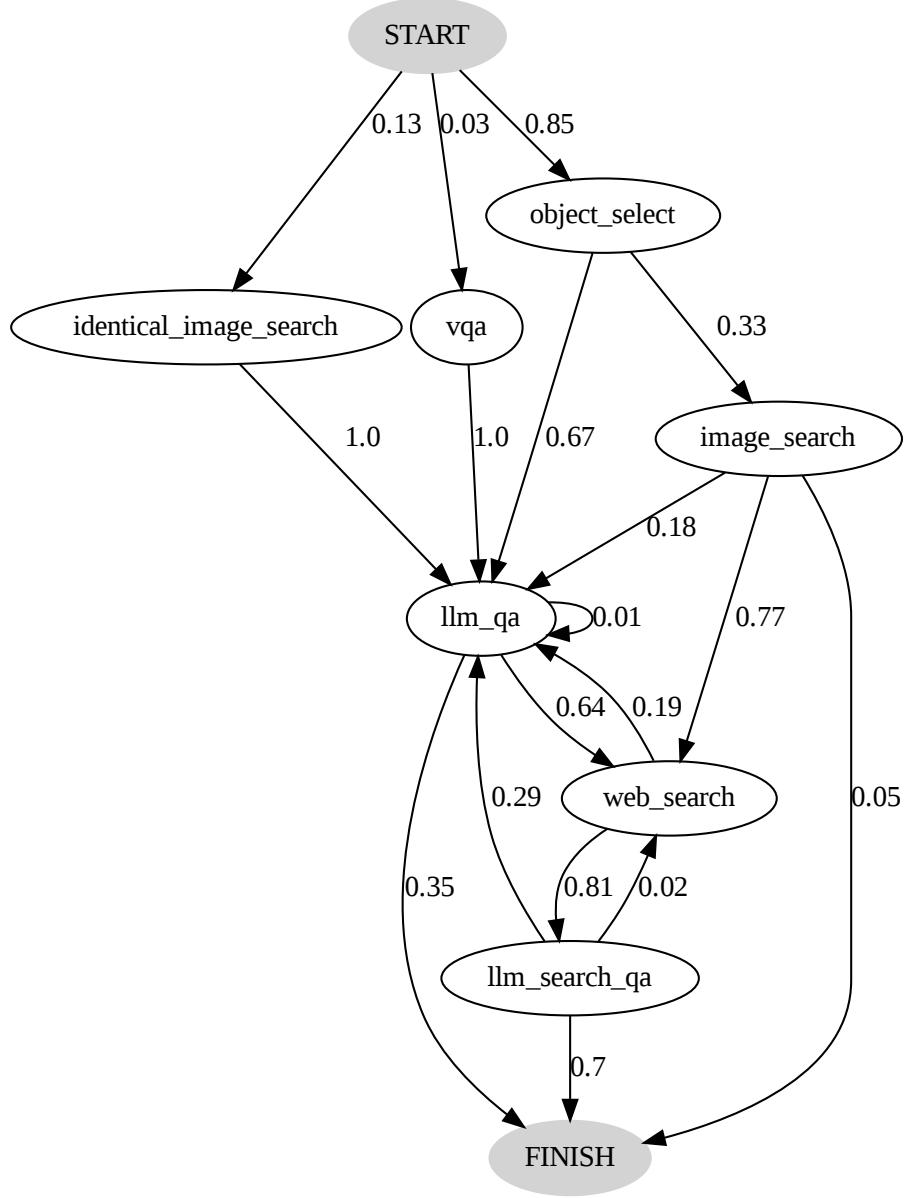


Figure 10: Induced transition frequency graph of AVIS over Infoseek dataset.

```

10 Query: %s
11 Context: %s
12 Action: \n
13 """
14
15 task_instructions = {
16 'vqa':
17     'You will ask simple question about this image to a external QA module. Please use this when the input
        query is very straightforward and simple.', \
18 'object_select':
19     'You will select one of the object we detect to dig further. Please use when the question asks about a
        specific object.', \
20 'identical_image_search':

```

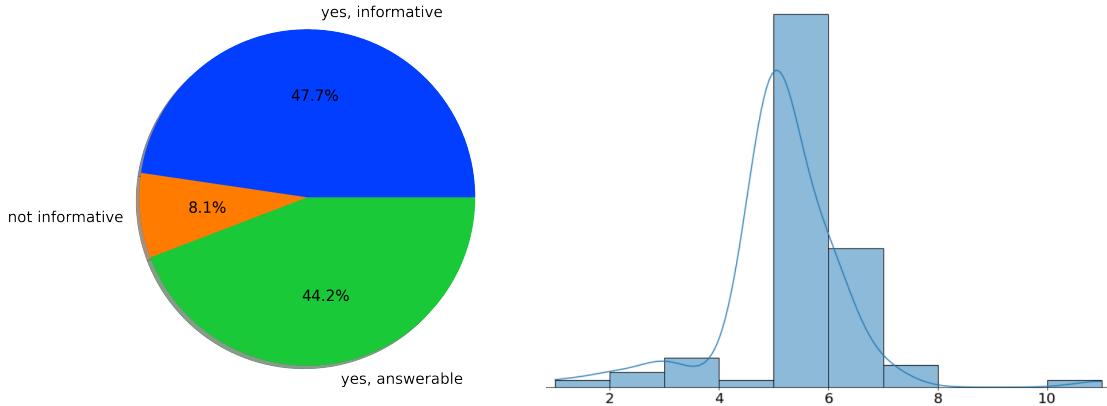


Figure 11: Overall frequency of judgement by reasoner of AVIS.

Figure 12: Length distribution of AVIS’s generated action sequences.

Model	Unseen Entity	Unseen Question
PALM (Q-only, few-shot)	6.6	4.8
OFA (fine-tune)	2.9	6.2
PALI (fine-tune)	5.9	13.3
PALM w/ CLIP (few-shot + external knowledge)	14.9	15.6
FiD w/ CLIP (fine-tune + external knowledge)	<u>17.6</u>	<u>18.9</u>
AVIS (ours, few-shot)	31.4	33.6

Table 5: **Visual Question Answering** results (accuracy) on Infoseek_{human}. The first four rows are results from their paper that do not use external knowledge, and the next two are from their paper that use CLIP as knowledge source.

```

21   'You will see captions of all images identical to the given image. Please use when the question asks
22   about the whole image instead of a part.', \
23   'You will see captions of all images similar to this object. Please use when you need more information.', \
24   'web_search':
25   'You will send question to Google Search to get knowledge. Please use when the current query requires
26   extra knowledge', \
27   'You will send question to a QA module. Please use this when the input query is simple and contain
28   common-sense knowledge'
}

```

Listing 1: Planner prompt skeleton and Task instructions

```

1 vqa_plan_prompts = [
2 """Query: what is the train carrying?
3 Context: [
4   a train traveling down train tracks next to a forest . There are four trains on the railway track. In the
      background there are trees,poles and sky. (Caption, whole image)
5   Extracted Text: BNSF (score=100.0),
6   BNSF Railway: BNSF Railway is one of the largest freight railroads in North America (score=89.3),
7 ]
8 Action: vqa
9 """
10 """Query: What is the girl wearing on her legs?
11 Context: [
12   a woman standing in a field putting on a coat . There is a woman standing on the ground. This is grass and
      there are plants. In the background we can see some trees and this is sky. (Caption, whole image)
13 ]
14 Action: vqa
15 """
16 """Query: what color is the bus?
17 Context: [
18   a double decker bus parked in front of a building . There is a double decker bus on the road and this is
      snow. Here we can see a pole, light, trees, and houses. In the background there is sky. (Caption, whole
      image)
19   Extracted Text: ENVIRO400 (score=100.0),
20   Extracted Text: Les Miserables (score=100.0),
21   Query Suggestion: les miserables (score=100.0),
22   Volvo Olympian: The Volvo Olympian was a rear-engined 2-axle and 3-axle double decker bus chassis
      manufactured by Volvo at its Irvine, Scotland factory (score=88.5),
23   Alexander Dennis Enviro400: The Alexander Dennis Enviro400 is a twin-axle low-floor double-decker bus that
      was built by the British bus manufacturer Alexander Dennis between 2005 and 2018 (score=85.4),
24 ]
25 Action: vqa
26 """
27 """Query: what is the person doing?
28 Context: [
29   two people sitting on the floor opening presents . There are sofas on the sofas there are pillows, here
      there is table, on the table there are plants and other objects, here there are two persons sitting on
      the ground, gift boxes, dog and this is floor. (Caption, whole image)
30 ]
31 Action: vqa
32 """
33 ]
34 object_select_plan_prompts = [
35 """Query: what is the name of this building?
36 Context: [
37   a group of people that are standing in front of a building . There is a building in the left corner which
      has few people standing in front of it and there is a fire hydrant in the right corner and there is a
      street light pole beside it. (Caption, whole image)
38   Query Suggestion: Alcatraz Warden's House San Francisco (score=100.0),
39   Alcatraz Island (historic_site): Alcatraz Island is a small island 1 (score=91.9),
40   Warden's House: The Warden's House was the home of the wardens of the federal penitentiary on Alcatraz
      Island, off San Francisco (score=78.1),
41 ]
42 Action: object_select
43 """
44 """Query: what is the island?
45 Context: [
46   a view of a mountain from a cable car . There is a ropeway. Behind that there are trees and hills.
      (Caption, whole image)
47   Ngong Ping 360 (gondola_lift_station): Ngong Ping 360 is a bicable gondola lift on Lantau Island in Hong
      Kong (score=91.8),
48   Tian Tan Buddha (monument): The Big Buddha is a large bronze statue of Buddha, completed in 1993, and
      located at Ngong Ping, Lantau Island, in Hong Kong (score=79.0),
49 ]
50 Action: object_select
51 """
52 """Query: what is the name of this place?
53 Context: [
54   a cemetery with a building in the background . There is a road and there are many atoms and trees beside it
      and there is a building in the right corner. (Caption, whole image)
55 ]
56 Action: object_select
57 """
58 """Query: what is the name of this bird?
59 Context: [
60   a bird sitting on top of a lush green hillside . There is a bird on the grassland in the foreground area of
      the image and the background is blurry. (Caption, whole image)
61   Atlantic puffin (type_of_bird): The Atlantic puffin, also known as the common puffin, is a species of
      seabird in the auk family (score=73.2),
62   Horned puffin (type_of_bird): The horned puffin is an auk found in the North Pacific Ocean, including the
      coasts of Alaska, Siberia and British Columbia (score=73.2),
63   Puffins (type_of_bird): Puffins are any of three species of small alcids in the bird genus Fratercula
      (score=73.2),
64   Fraterculini (score=48.8),
65   Auk (type_of_bird): An auk or alcid is a bird of the family Alcidae in the order Charadriiformes
      (score=11.8),
66 ]
67 Action: object_select
68 """
69 ]
70 identical_image_search_plan_prompts = [
71 """Query: what is the name of this building?
72 Context: [

```

```

73 a group of people that are standing in front of a building . There is a building in the left corner which
    has few people standing in front of it and there is a fire hydrant in the right corner and there is a
    street light pole beside it. (Caption, whole image)
74 Query Suggestion: Alcatraz Warden's House San Francisco (score=100.0),
75 Alcatraz Island (historic_site): Alcatraz Island is a small island 1 (score=91.9),
76 Warden's House: The Warden\`s House was the home of the wardens of the federal penitentiary on Alcatraz
    Island, off San Francisco (score=78.1),
77 ]
78Action: identical_image_search
79 """
80 """Query: what is the aircraft?
81Context: [
82 a fighter jet sitting on top of an airport tarmac . There is a plane and missiles on the ground. At the
    left a person is standing wearing a cap. (Caption, whole image)
83 Extracted Text: AIRLINERS.NET (score=100.0),
84 Query Suggestion: airliners.net (score=100.0),
85 Airliners: Airliners (score=74.8),
86 British Aerospace Hawk 200: The British Aerospace Hawk 200 is a single-seat, single engine light multirole
    fighter designed for air defence, air denial, anti-shipping, interdiction, close air support, and
    ground attack (score=74.8),
87 product: Airfix BAE Hawk T1 1:72 (score=0.0),
88 product: Rolls-royce Adour In The Hawk / Bae Hawk 200 . Pdf/download (score=0.0),
89 ]
90Action: identical_image_search
91 """
92 """Query: what is the name of this place?
93Context: [
94 a row of pillars sitting next to a dirt road . There is a building and this is plant. Here we can see
    pillars and a sky. (Caption, whole image)
95 Query Suggestion: Palmyra Archaeology (score=100.0),
96 Great Colonnade at Palmyra (ancient_roman_architecture_structure): The Great Colonnade at Palmyra was the
    main colonnaded avenue in the ancient city of Palmyra in the Syrian Desert (score=90.3),
97 ]
98Action: identical_image_search
99 """
100 """Query: what is the name of this lake?
101Context: [
102 a view of a river surrounded by mountains . There are trees in the right corner and there is a river and
    mountains in front of it. (Caption, whole image)
103 Monte Bre (peak): Monte Bre is a small mountain east of Lugano on the flank of Monte Boglia with a view of
    the bay of Lugano and the Pennine Alps and the Bernese Alps (score=85.5),
104 product: Top Searcher (score=0.0),
105 ]
106Action: identical_image_search
107 """
108 ]
109action_prompt_dict = {'vqa': vqa_plan_prompts, 'object_select': object_select_plan_prompts,
    'identical_image_search': identical_image_search_plan_prompts, 'image_search':
        image_search_plan_prompts, 'web_search': web_search_plan_prompts,
    'llm_qa': llm_qa_plan_prompts}
110

```

Listing 2: Planning Prompts Example

We then show how AVIS decompose question into a visual sub-question and a knowledge sub-question. This is done at beginning to guide later tool usage.

```

1 question_decomposition_prompt = """
2     Read the following question for a given image. Decompose the question into two sub-questions.
3
4     The first will ask information about the image, and the second requires reasoning over the textual
        knowledge.
5     In the second question, we use # to denote the answer of the first question.
6
7
8     Question: what chemical makes the vegetable orange?
9     Visual: which orange vegetable is shown?
10    Knowledge: chemical makes # orange?
11
12
13    Question: How long can their horns grow?
14    Visual: which animals are shown?
15    Knowledge: How long can #'s horns grow?
16
17
18    Question: What is a competition for these animals called?
19    Visual: which animals are shown?
20    Knowledge: competition for #?
21
22
23    Question: What is the name of the ancient greek sport that evolved into the sport featured above?
24    Visual: which sport is played?
25    Knowledge: name of the ancient greek sport that evolved into #?
26
27
28    Question: Which food item here has the most protein?
29    Visual: what are the food items shown?
30    Knowledge: Which food item of # has the most protein?
31
32
33    Question: How many calories are in this meal?

```

```

34 Visual: what are the food items shown?
35 Knowledge: calories in #?
36
37
38 Question: What type of sandwich is this?
39 Visual: which type of sandwich is shown?
40 Knowledge: #
41
42 Question: What is the name of the restaurant where this was served?
43 Visual: which food items are served?
44 Knowledge: restaurant where # was served?
45
46
47 Question: What genus of bird is flying here?
48 Visual: what genus of bird is flying?
49 Knowledge: #
50
51
52 Question: What is the main ingredient in this food?
53 Visual: which food is shown?
54 Knowledge: main ingredient in #?
55 """

```

Listing 3: Question Decomposition Prompts

Below are several examples to help AVIS learns how to select the most suitable object #ID.

```

1 object_select_prompt = """
2     Please think step by step. In the following, you will be given a "Query", a list of "Objects".
3
4     Your task is to predict the object #ID that is mostly relevant to answer the querys. Please generate the
5         detailed explanation why you select this object, and then output ID in "Object #ID".
6
7 Query: which city is this place?
8 Object #0 [
9     a row of pillars sitting next to a dirt road . There is a building and this is plant. Here we can see
10        pillars and a sky. (Caption, whole image)
11     Query Suggestion: Palmyra Archaeology (score=100.0),
12     Great Colonnade at Palmyra (ancient_roman_architecture_structure): The Great Colonnade at Palmyra was the
13         main colonnaded avenue in the ancient city of Palmyra in the Syrian Desert (score=90.3),
14 ]
13 Object #1 [
14     a green plant sitting next to a brick wall . There is a plant and this is wall. And there is a sky.
15        (Caption, center)
15     Date palm (type_of_palm_trees): Phoenix dactylifera, commonly known as date palm, is a flowering plant
16         species in the palm family, Arecaceae, cultivated for its edible sweet fruit called dates (score=81.7),
17 ]
17 Object #2 [
18     a wicker basket sitting on top of a rock . There is a blur image of a rock. (Caption, lower right)
19 ]
20 Output: The query asks about the city of the place. Only Object #0 (whole image) mentions city name Palmyra,
21         which is an ancient city. Also, Object #0 contains Query Suggestion "Palmyra Archaeology".
21 Therefore, the predicted Object #ID is 0.
22
23
24 Query: where is this place?
25 Object #0 [
26     a view of a valley surrounded by mountains . There are hills and this is grass. Here we can see trees and
27         this is sky. (Caption, whole image)
27 ]
28 Object #1 [
29     a view of a lush green hillside with trees . There is a house on the rock and there are few plants beside
30         it and there is a greenery ground in the background. (Caption, center)
30     Monterey Pine (type_of_conifers): Pinus radiata, the Monterey pine, insignis pine or radiata pine, is a
31         species of pine native to the Central Coast of California and Mexico (score=49.1),
31     European rabbit (type_of_leporids): The European rabbit or coney is a species of rabbit native to the
32         Iberian Peninsula, western France, and the northern Atlas Mountains in northwest Africa (score=31.3),
32 ]
33 Object #2 [
34     a green plant growing on a rocky surface . There is a blur image of trees and rocks. (Caption, lower center)
35     product: GreenView Fairway Formula Seed Success Pailles biodegradable avec engrais Sac de 4.5 kg Couvre 200
36         m2 (score=0.0),
36 ]
37 Object #3 [
38     a rocky hillside with lots of green vegetation . There are trees and this is rock. (Caption, lower left)
39     Willow: Willows, also called sallows and osiers, of the genus Salix, comprise around 350 species of
40         typically deciduous trees and shrubs, found primarily on moist soils in cold and temperate regions
40         (score=31.3),
40     Tamarisk: The genus Tamarix is composed of about 50-60 species of flowering plants in the family
41         Tamaricaceae, native to drier areas of Eurasia and Africa (score=26.8),
41 ]
42 Output: The query asks about the location of this place. Although these entries doesn't explicitly contain
42         location name, but Object #1 (center) contains Monterey Pine and European rabbit, which might hint the
42         location later.
43 Therefore, the predicted Object #ID is 1.
44 """

```

Listing 4: Object Select Prompts

Below are the prompts to extract answer from objects and extracted captions of similar images.

```
1 reason_vqa_prompt = """
2 Please think step by step. In the following, you will be given:
3
4- Query: The query to be asked.
5- Think: Why the following knowledge is retrieved.
6- Entity: A list of entities that describe the object.
7- Retrievals: A list of web documents that are similar to the object. If there's "high confidence", it's very
     important.
8
9 Your task is to predict a short answer to the query based on the provided information. You need to first
     identify which knowledge entry is mostly relevant, and then extract the answer from the knowledge.
10 Rely on Object information more, and if there contains "Query Suggestion", try to use it. Otherwise, if a
     information appears lots of time, there's a higher chance it's the answer.
11 After explaining your decision choice, saying "Answer is" and appending your predicted short answer. Please
     also generate the type of the answer after a comma.
12 If you are uncertain about the answer, especially when the knowledge is irrelevant to the query, say "cannot
     be answered". Do not generate the answer not inside the provided knowledge.
13
14
15
16 Query: what is this building?
17 Think: object (whole image) contains stockholm city hall, which is the seat of stockholm municipality in
     stockholm, sweden.
18 Object: [
19   Stockholm City Hall (city_hall): Stockholm City Hall is the seat of Stockholm Municipality in Stockholm,
     Sweden (score=96.1),
20   Bla Hallen (banquet_hall): The Blue Hall is the main hall of the Stockholm City Hall best known as the
     banquet hall for the annual Nobel Banquet, and also used for state visits, student balls, jubilees and
     other large events (score=79.0),
21 ]
22 Retrievals: [
23   high confidence: City Hall - Blue Hall (1) | Stockholm (2) | Pictures | Sweden in Global-Geography
     (score=47.8),
24   high confidence: le salon bleu a city hall (salle de remise des prix nobel) - Picture of Stockholm,
     Stockholm County - Tripadvisor (score=47.7),
25 ]
26
27 Output: The query asks about the building. From both Object and Retrievals, there are mentions about
     Stockholm City Hall and Blue Hall. As Stockholm City Hall contains Blue Hall, the answer shall be
     Stockholm City Hall.
28 Therefore, the predicted answer is Stockholm City Hall.
29
30
31 Query: which sport is played?
32 Think: Object shows a snail sitting on top of a tennis ball.
33 Object: [
34   Cantareus apertus (type_of_gastropods): Cantareus apertus, commonly known as the green garden snail, is a
     species of air-breathing land snail, a terrestrial pulmonate gastropod mollusc in the family Helicidae,
     the typical snails,
35   Garden snail (type_of_gastropods): Cornu aspersum, known by the common name garden snail, is a species of
     land snail in the family Helicidae, which includes some of the most familiar land snails,
36   Helix aspersa aspersa (type_of_gastropods),
37   Slug: Slug, or land slug, is a common name for any apparently shell-less terrestrial gastropod mollusc,
38   Snail: A snail is a shelled gastropod,
39 ]
40 Retrievals: [
41   2019 NEWBIE Competition Winner Steven Ryan, Snail Farming - YouTube,
42   Alive specimens. a. Megalobulimus ovatus (CMIOC 11136), b. Thaumastus... | Download Scientific Diagram,
43   Brown garden snail Manaaki Whenua,
44   Common garden snail and baby,
45   Easy Everyday Food for Garden Snails - Ask the plantician,
46   Green Life Soil: Natural pest & disease control in a winter garden,
47   Helminthoglyptinae - Wikipedia,
48   Hydrosalpingitis in broilers - Veterinaria Digital,
49   Master Gardener: Protecting squash and cucumbers from slugs and snails - Press Enterprise,
50   Mother Baby Blue Snails On Phalaenopsis Stock Photo 530400856 | Shutterstock,
51 ]
52
53 Output: The query asks about sport. From both entities and retrievals, they only talk about snail, and there
     is no information about which sport is played.
54 Therefore, given the provided information, this query cannot be answered.
55
56
57 Query: which sport is played?
58 Think: object , object , and object all contain people playing basketball. however, object is the only one
     that contains a group of women playing basketball.
59 therefore, the predicted object #id is 0.
60 Retrievals: [
61   08.07.2011 Zanele Mdodana of South Africa in action during the Quarter-finals between New Zealand and South
     Africa, Mission Foods World Netball Championships 2011 from the Singapore Indoor Stadium in Singapore
     Stock Photo - Alamy,
62   55 Brazilian Handball Team Images, Stock Photos & Vectors | Shutterstock,
63   ::::Malawi High Commission::::,
64   Amanda Mynhardt Photostream | Netball, Netball singapore, Netball south africa,
65   Australia pass Malawi test with flying colours at Netball World Cup | Netball World Cup 2019 | The Guardian,
66   Australia's Jo Weston (second left) and Barbados' Latonia Blackman in action during the Netball World Cup
     match at the M&S Bank Arena, Liverpool Stock Photo - Alamy,
67   Birmingham 29795 World Netball Championships Final Editorial Stock Photo - Stock Image | Shutterstock,
68   Bridget kumwenda malawi netball hi-res stock photography and images - Alamy,
```

```

69 England V Australia International Netball Series Photos and Premium High Res Pictures | Netball, Netball
    quotes, Inspirational women,
70 File:Xx0992 - Madrid basketball Donna Burns - 3b - Scan.jpg - Wikimedia Commons,
71 ]
72
73Output: The query asks about which sport is played. From retrievals, there exist many mentions about netball,
    and mentions that they are played by women.
74therefore, the predicted answer is women netball.
75
76
77
78Query: what is the name of the insect?
79Think: only object (while image) mentions the name of the insect, western tiger swallowtail.
80Object: [
81 Query Suggestion: Western Tiger Swallowtail (score=100.0),
82 Canadian tiger swallowtail (type_of_lepidoptera): Papilio canadensis, the Canadian tiger swallowtail, is a
    species of butterfly in the family Papilionidae (score=78.4),
83 Eastern tiger swallowtail (us_state_butterfly): Papilio glaucus, the eastern tiger swallowtail, is a
    species of butterfly native to eastern North America (score=78.4),
84 ]
85Retrievals: [
86 high confidence: kupu-kupu - Wiktionary (score=100.0),
87 high confidence: Top Spots for Nature Watching and Birding | VisitMaryland.org (score=100.0),
88 high confidence: File:Eastern Tiger Swallowtail Papilio glaucus on Milkweed 2800px.jpg - Wikimedia Commons
    (score=99.8),
89 high confidence: Photographing Butterflies - Life in the Finger Lakes (score=97.8),
90 ]
91
92Output: The query asks about the name of the insect. From Object, it contains a very informative "Query
    Suggestion: Western Tiger Swallowtail".
93Therefore, the predicted answer is Western Tiger Swallowtail.
94
95"""

```

Listing 5: Reason Prompt (Visual Question)

Below are prompts AVIS extract answer from search results:

```

1reason_qa_prompt = """
2    Please think step by step. In the following, you will be given a "Query", and a list of "Knowledge" from
        Google Search related to this query.
3
4    Your task is to predict a short answer to the query based on the provided information. You need to first
        identify the most relevant knowledge entry, and then predict a short answer based on the knowledge. If
        a information appears lots of time, there's a higher chance it's the answer.
5
6    After explaining your decision choice, saying "Answer is" and appending your predicted answer.
7    If you are uncertain about the answer, especially when the knowledge is irrelevant to the query, say
        "cannot be answered". Do not generate the answer not inside the provided knowledge.
8
9
10Query: What chemical makes carrot orange?
11Knowledge: [
12Title: How did carrots become orange? - The Economist
13Content: High Confidence Response: carotenoids.
14
15Context: The chemical compounds that give carrots their vivid colour, carotenoids, are usually used by plants
        that grow above ground to assist in the process of photosynthesis.
16
17Title:
18Content: carotenoids
19
20The chemical compounds that give carrots their vivid colour, carotenoids, are usually used by plants that
        grow above ground to assist in the process of photosynthesis.
21
22Title: Can Eating Too Many Carrots Make Your Skin Turn Orange? | Britannica - Encyclopedia Britannica
23Content: Maybe not! Carrots and other orange fruits and vegetables are rich in a pigment known as
        beta-carotene. In humans, this pigment is converted to vitamin A by specialized cells in the small
        intestine. When high levels of beta-carotene are consumed, not all of the pigment is converted to
        vitamin A.
24Fortunately, the skin discoloration fades when the diet is changed and the levels of beta-carotene in the
        blood decline.
25
26]
27Output: The query asks about chemical that makes carrot orange. Because there's one high confidence result,
        the most relevant knowledge entries about such chemical is "High Confidence Response: carotenoids."
28From this result we know the chemical shall be carotene.
29Therefore, the predicted answer is carotene.
30
31
32
33Query: What is the name of the drainage basin of ounasjoki?
34Knowledge: [
35Title: Ounasjoki - Wikipedia
36Content: It is also the largest river entirely within its borders. Ounasjoki is approximately 299.6
        kilometres (186.2 mi) in length, and the catchment area is 13,968 square kilometres (5,393 sq mi), 27%
        of the Kemijoki catchment area.
37Tributaries
38
39 - Nakkalajoki.
40 - Kakkalojoki.

```

41 - Syva Tepastojoki.
 42 - Loukinen.
 43 - Meltausjoki.
 44 Course. The Ounasjoki originates at Ounasjarvi lake in Enontekio. It flows first eastwards through Perilajarvi lake and turns south after some seven kilometres. The river then follows southern-sou
 45
 46 Title: DRAINAGE BASIN OF THE BALTIC SEA - UNECE
 47 Content: Vistula. 194,424. Baltic Sea. BY, PL, SK, UA. - Bug. 39,400. Vistula. BY, PL, UA. - Dunajec. 4726.7. Vistula. PL, SK. -Poprad. 2,077. Dunajec. PL, SK. Oder. 118,861. Baltic Sea. CZ, DE, PL. - Neisse Oder. CZ, DE, PL. - Olse Oder. CZ, PL. 1 The assessment of water bodies in italics was not included in the present publication. 2 For the Venta River Basin District, which includes the basins of the Barta/Bartuva and Sventoji rivers. Oulu. Lulea. Rovaniemi. Lake. Oulujarvi. Lake. Tornetrask. Torne. Oulujoki.
 48]
 49 Output: The query asks about drainage basin of ounasjoki. The most relevant knowledge entry that contain basin is "Venta River Basin District, which includes the basins of the Barta/Bartuva and Sventoji rivers."
 50 From this result we know the drainage basin shall be Venta River Basin.
 51 Therefore, the predicted answer is Venta River Basin.
 52
 53
 54 Query: What is the typical diameter (in centimetre) of tennis?
 55 Knowledge: [
 56 Title: What Size Is A Tennis Ball In Cm? - Metro League
 57 Content: To Recap. A tennis ball is typically about 2 cm in diameter. Similar Posts: What Is A Junk Ball In Tennis?
 58 How tall is a tennis ball? Tennis Balls come in different sizes, some as small as 2.575"-2.7" (6.54-6.86 cm) and others up to 8 inches (20 cm). The mass of a tennis ball must be between 1.975-2.095 oz (56-59 g).
 59
 60 Title: Tennis Ball Dimensions & Drawings | Dimensions.com
 61 Content: Tennis Balls have a diameter of 2.575"-2.7" (6.54-6.86 cm) and circumference of 8.09"-8.48" (20.6-21.5 cm). The mass of a Tennis Ball must be between 1.975-2.095 oz (56-59.4 g).
 62 Tennis Balls have a diameter of 2.575"-2.7" (6.54-6.86 cm) and circumference of 8.09"-8.48" (20.6-21.5 cm). The mass of a Tennis Ball must be between 1.975-2.095 oz (56-59.4 g). A Tennis Ball is a ball designed for the sport of tennis.
 63
 64 Title: Tennis ball - Wikipedia
 65 Content: Modern tennis balls must conform to certain criteria for size, weight, deformation, and bounce criteria to be approved for regulation play. The International Tennis Federation (ITF) defines the official diameter as 6.54-6.86 cm (2.57-2.70 inches). Balls must have masses in the range 56.0-59.4 g (1.98-2.10 ounces).
 66]
 67 Output: The query asks about diameter of tennis (in centimetre). the most relevant knowledge entry about diameter of tennis is "tennis balls have a diameter of 2.575"-2.7" (6.54-6.86 cm) and circumference of 8.09"-8.48" (20.6-21.5 cm)".
 68 As the query ask about centimetre, cm. From this result we know the diameter shall be 6.54 - 6.86.
 69 Therefore, the predicted answer is 6.54 - 6.86.
 70
 71
 72
 73 Query: Who is the inventor of women netball, sport?
 74 Knowledge: [
 75 Title:
 76 Content: History of netball - Wikipedia
 77
 78 In 1893, Martina Bergman-osterberg informally introduced one version of basketball to her female physical training students at the Hampstead Physical Training College in London, after having seen the game being played in the United States.
 79
 80 Title: History of netball - Wikipedia
 81 Content: In 1893, Martina Bergman-osterberg informally introduced one version of basketball to her female physical training students at the Hampstead Physical Training College in London, after having seen the game being played in the United States. Madame osterberg advocated physical fitness for women to better prepare them for motherhood and in the wider context of women's emancipation.
 82
 83 Title: Netball - Wikipedia
 84 Content: A common misunderstanding of netball's origins has resulted in the mistaken belief that netball was created to prevent women from playing basketball. However, netball's development traces back to American sports teacher Clara Gregory Baer's misinterpretation of the basketball rule book in 1895.
 85 History. Netball's early development emerged from Clara Baer's misinterpretation of the early rules of James Naismith's new sport of basketball (which he developed while studying in Massachusetts) and eventually evol
 86
 87 Title: History of Netball - World Netball
 88 Content: Women's indoor basketball began exactly two days later when female teachers to the gym were captivated by the game but it wasn't until 1895 that the current game of netball was well and truly shaped. When Clara Baer, a sports teacher in New Orleans, wrote to Naismith asking for a copy of the rules, the subsequent rules package contained a drawing of
 89]
 90 Output: The query asks about inventor of women netball. The most relevant knowledge entry about women netball inventor is "In 1893, Martina Bergman-Osterberg informally introduced one version of basketball to her female physical training students".
 91 From the result, we know the inventor shall be Martina Bergman-Osterberg.
 92 Therefore, the predicted answer is Martina Bergman-Osterberg.
 93
 94
 95 Query: How many elevators does torre picasso have?
 96 Knowledge: [
 97 Title:
 98 Content: Torre Picasso | Turismo Madrid
 99

```

100 The interior of the Picasso Tower houses offices designed as intelligent spaces equipped with the highest
     technology, comfort and use of space. It has 18 lifts, divided into three groups of six.
101
102 Title: Torre Picasso - Wikipedia
103 Content: 26 elevators; 18 serve office floors divided into three zones:
104
105 - 1st-18th floors at 2.5 m/s (8.20 ft/s)
106 - 18th-32nd floors at 4 m/s (13.12 ft/s)
107 - 32nd-43rd floors at 6 m/s (19.69 ft/s) (fastest in Spain)
108
109 Title: Torre Picasso - Field Trip
110 Content: 26 elevators, of which 18 to office floors in 3 groups of 6:
111
112 - 1st-18th floors at 2.5 m/s (8.20 ft/s)
113 - 18th-32nd floors at 4 m/s (13.12 ft/s)
114 - 32nd-43rd floors at 6 m/s (19.69 ft/s) (apparently fastest in Spain)
115
116 Title: Torre Picasso - Wikiwand
117 Content: The building as seen from the junction of the Paseo de la Castellana and the Plaza de Pablo Ruiz
     Picasso. 26 elevators; 18 serve office floors divided into three zones: 1st-18th floors at 2.5 m/s
     (8.20 ft/s) 18th-32nd floors at 4 m/s (13.12 ft/s)
118
119 ]
120 Output: The query asks about number of elevators in torre picasso. the most relevant knowledge entry about
     number of elevators in torre picasso is "26 elevators; 18 serve office floors divided into three
     zones:".
121 From the result, we know the number of elevators shall be 26.
122 therefore, the predicted answer is 26.
123 """

```

Listing 6: Reason Prompt (Knowledge Question)

```

1 class MemoryState:
2     state: str = ''
3     traversed_actions: list = []
4     query: str = ''
5     context: str = ''
6
7     def __init__(self, state, query = '', context = ''):
8         self.state = state
9         self.query = query
10        self.context = context
11
12    def plan(transition_graph, cur_memory, lens_res, retr_res):
13        action_list = [a for a in transition_graph[cur_memory.state] if a not in cur_memory.traversed_actions]
14        action_prompt = ''
15        for a in action_list:
16            action_prompt += ' --> ' + a + ': ' + task_instructions[a] + '\n'
17        prompt_example = ""
18        for a in action_list:
19            prompt_example += action_prompt_dict[a] + "\n"
20        action_prompt = planner_prompt % (cur_memory.query, action_prompt, prompt_example, cur_memory.query,
21                                         cur_memory.context)
22        action = api_utils.call_palm(action_prompt)[0]
23
24        instruction = []
25        if action in require_instruction:
26            exclude_ids = cur_memory.traversed_actions:
27            prompt = instruction_prompt(cur_memory.query, lens_res, exclude_ids)
28            res = api_utils.call_palm(prompt)[0]
29            reason = parse_reason('the query asks about ' + reason)
30            instruction = [reason, res]
31        return action, instruction
32
33    def avis_execution(d):
34        state = 'START'
35        answer = None
36
37        prompt = question_decomposition_prompt + 'Question: ' + q + '\n'
38        res = api_utils.call_palm(prompt)[0]
39
40        vqi = res.find('Visual: ')
41        kqi = res.find('Knowledge: ')
42        vq = res[vqi + 8: kqi-1]
43        kq = res[kqi+1:]
44
45        working_memory = [MemoryState(state = 'START', query = vq, context = lens_res[0])]
46        while not answer:
47            cur_memory = working_memory[-1]
48            action, instruction = plan(transition_graph, cur_memory, lens_res, retr_res)
49            exec_res = execute(action, instruction, lens_res, retr_res)
50            res = reason(exec_res)
51            if 'not informative' in res:
52                cur_memory.traversed_actions += [action]
53            elif 'answer is' in res:
54                answer = res[10:]
55            else:
56                working_memory += [MemoryState(state = action, query = kq, context = res)]
57
58    return answer

```

Listing 7: Workflow of AVIS (code snippets)